

Machine Learning 1

Lecture 7.3 - Supervised Learning
Classification - Logistic Regression:
Stochastic Gradient Descent

Erik Bekkers

(Bishop 4.3.2)



Logistic Regression for Two Classes

- ▶ Given: Dataset $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_N)^T$ with binary targets $\mathbf{t} = (t_1, \dots, t_N)^T$ with $t_n \in \{\mathcal{C}_1, \mathcal{C}_2\} = \{1, 0\}$

- ▶ Conditional likelihood function:

$$p(\mathbf{t}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(t_n|\mathbf{x}_n, \mathbf{w}) = \prod_{n=1}^N y_n^{t_n} (1 - y_n)^{1-t_n}$$

$$y_n = p(C_1|\phi_n) = \sigma(w^T \phi_n) \quad \phi_n = \phi(\mathbf{x}_n)$$

- ▶ Maximizing the conditional likelihood/minimizing the cross-entropy

$$E(\mathbf{w}) = -\ln p(\mathbf{t}, \mathbf{X}, \mathbf{w}) = -\sum_{n=1}^N t_n \ln y_n + (1 - t_n) \ln(1 - y_n)$$

- ▶ $E(\mathbf{w})$: convex, but no closed form solution!

$$y_n = \sigma(\mathbf{w}^T \phi_n) \text{ is nonlinear in } \mathbf{w}$$

Logistic Regression (K=2): SGD

$$y_n = \sigma(\mathbf{w}^T \phi_n)$$

- Stochastic Gradient Descent for cross-entropy:

$$E(\mathbf{w}) = - \sum_{n=1}^N t_n \ln y_n + (1 - t_n) \ln(1 - y_n) = \sum_{n=1}^N E_n(\mathbf{w})$$

- Update rule given a random data point (\mathbf{x}_n, t_n)

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_n(\mathbf{w}^{(\tau)})^T$$

- Gradient: $\nabla E_n(\mathbf{w}) = \left(\frac{\partial E_n(\mathbf{w})}{\partial w_0}, \dots, \frac{\partial E_n(\mathbf{w})}{\partial w_{M-1}} \right)$

$$\frac{\partial E_n(\mathbf{w})}{\partial w_j} = \frac{\partial E_n(\mathbf{w})}{\partial y_n} \frac{\partial y_n}{\partial w_j} = \left(-\frac{t_n}{y_n} + \frac{1-t_n}{1-y_n} \right) \cdot \frac{\partial y_n}{\partial w_j}$$

$$\frac{\partial y_n}{\partial w_j} = \frac{\partial}{\partial w_j} \sigma(\mathbf{w}^T \phi_n)$$

Logistic Regression (K=2): SGD

$$\triangleright \frac{\partial y_n}{\partial w_j} = \frac{\partial}{\partial w_j} \sigma(\mathbf{w}^T \phi_n)$$

$$\frac{\partial \mathbf{w}^T \phi_n}{\partial w_j} = \frac{\partial}{\partial w_j} \sum_{i=0}^{M-1} w_i \phi_{ni} = \phi_{nj}$$

$$\triangleright \text{Use } \frac{\partial \sigma(a)}{\partial a} = \sigma(a)(1 - \sigma(a))$$

$$\triangleright \frac{\partial}{\partial w_j} \sigma(\mathbf{w}^T \phi_n) = \overbrace{\sigma(\mathbf{w}^T \phi_n)}^{y_n} (1 - \overbrace{\sigma(\mathbf{w}^T \phi_n)}^{y_n}) \cdot \frac{\partial \mathbf{w}^T \phi_n}{\partial w_j}$$

$$= y_n (1 - y_n) \phi_{nj}$$

$$\triangleright \frac{\partial E_n(\mathbf{w})}{\partial w_j} = - \frac{t_n}{y_n} \frac{\partial y_n}{\partial w_j} + \frac{1 - t_n}{1 - y_n} \frac{\partial y_n}{\partial w_j}$$

$$= - \frac{t_n}{y_n} \cdot y_n (1 - y_n) \phi_{nj} + \frac{1 - t_n}{1 - y_n} \cdot y_n (1 - y_n) \phi_{nj}$$

$$= -t_n \phi_{nj} + t_n y_n \phi_{nj} + y_n \phi_{nj} - t_n y_n \phi_{nj} = (y_n - t_n) \phi_{nj}$$

(activation fn
= link fn
Bishop 4.3.6)

Logistic Regression (K=2): SGD

- ▶ Stochastic Gradient Descent for cross-entropy:

$$E(\mathbf{w}) = - \sum_{n=1}^N t_n \ln y_n + (1 - t_n) \ln(1 - y_n) = \sum_{n=1}^N E_n(\mathbf{w})$$

- ▶ Update rule given a random data point (\mathbf{x}_n, t_n)

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_n(\mathbf{w}^{(\tau)})^T$$

- ▶ $\frac{\partial E_n(\mathbf{w})}{\partial w_j} = (y_n - t_n) \phi_j(\mathbf{x}_n)$

- ▶ Gradient: $\nabla E_n(\mathbf{w})^T = \left(\frac{\partial E_n(\mathbf{w})}{\partial w_0}, \dots, \frac{\partial E_n(\mathbf{w})}{\partial w_{M-1}} \right)^T = (y_n - t_n) \phi_n$

- ▶ Update rule:

$$\mathbf{w}^{(\tau+1)} = \underline{\mathbf{w}}^{(\tau)} - \eta \underline{(y_n - t_n)} \underline{\phi_n}$$

perceptron

Stochastic Gradient Descent

1. Initialize $\mathbf{w}^{(0)}$
2. Choose a learning rate η
3. While $\|\mathbf{w}^{(\tau+1)} - \mathbf{w}^{(\tau)}\| > \varepsilon$

I. Choose a random data point (\mathbf{x}_n, t_n)

II. Update \mathbf{w} :

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta (y_n^{(\tau)} - t_n) \phi(\mathbf{x}_n)$$

- ▶ If η too large: no convergence
- ▶ If η too small: very slow convergence
- ▶ Converged \mathbf{w}^* : estimate of minimizer of $E(\mathbf{w})$!

