


Machine Learning 1

Lecture 8.1 - Supervised Learning
Neural Networks

Erik Bekkers

(Bishop 5.1)



Fixed Basis Functions

Dataset: inputs $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_N)^T$ and targets $\mathbf{t} = (t_1, \dots, t_N)^T$

$$\underline{x}_n \in \mathbb{R}^D$$

Previously:

► Fixed features: $\phi(\mathbf{x}) = (\phi_0(\mathbf{x}), \dots, \phi_M(\mathbf{x}))^T$, $\phi_0(\mathbf{x}) = 1$

► Linear regression: $y(\mathbf{x}, \mathbf{w}) = \underline{\mathbf{w}}^T \underline{\phi(\underline{x})}$ $t_n \in \mathbb{R}$

► Classification: $y(\mathbf{x}, \mathbf{w}) = f(\underline{\mathbf{w}^T \phi(\mathbf{x})})$ $t_n \in \{0, 1\}$

f : nonlinear activation function e.g. (logistic sigmoid)

Neural Networks: Adaptive Basis Functions

Dataset: inputs $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_N)^T$ and targets $\mathbf{t} = (t_1, \dots, t_N)^T$

$$\underline{x}_n = (1, x_{n1}, \dots, x_{nd})^T \in \mathbb{R}^{D+1}$$

Neural networks:

- ▶ Create flexible non-linear features and learn them!

$$\phi_m(\mathbf{x}, \mathbf{w}_m^{(1)}) = h((\mathbf{w}_m^{(1)})^T \mathbf{x}) = h\left(\sum_{d=0}^D w_{md}^{(1)} x_d\right)$$

non-linear
activation fn

linear

$$\mathbf{W}^{(1)} = \begin{pmatrix} 1 & w_{11}^{(1)} & w_{12}^{(1)} & \dots & w_{1D}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} & w_{23}^{(1)} & \dots & w_{2D}^{(1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ w_{M1}^{(1)} & w_{M2}^{(1)} & w_{M3}^{(1)} & \dots & w_{MD}^{(1)} \end{pmatrix}$$

- ▶ Regression:

$$y(\mathbf{x}, \mathbf{W}^{(1)}, \mathbf{w}^{(2)}) = \sum_{m=0}^M w_m^{(2)} h\left(\sum_{d=0}^D w_{md}^{(1)} x_d\right) = \underline{w}^{(2)T} \underbrace{h(\mathbf{W}^{(1)T} \underline{x})}_{\phi(\underline{x})}$$

- ▶ Classification:

$$y(\mathbf{x}, \mathbf{W}^{(1)}, \mathbf{w}^{(2)}) = \underset{\sigma, \text{ soft max}}{f}\left(\underline{w}^{(2)T} \underbrace{h(\mathbf{W}^{(1)T} \underline{x})}_{\phi(\underline{x})}\right)$$

$$\underline{\phi}(\underline{x})$$


2-layer NN

Multilayer Perceptron (MLP): 2 layers

Model:

$$y_k(\mathbf{x}, \mathbf{W}^{(1)}, \mathbf{W}^{(2)}) = h^{(2)} \left(\sum_{m=0}^M w_{km}^{(2)} h^{(1)} \left(\underbrace{\sum_{d=0}^D w_{md}^{(1)} x_d}_{a_m} \right) \right)$$

Network diagram:



Input units x_d

$$\underline{x} = \begin{pmatrix} x_1 \\ \vdots \\ x_D \end{pmatrix} \in \mathbb{R}^{D+1}$$

Activations a_m

$$\underline{a} = \begin{pmatrix} a_1 \\ \vdots \\ a_M \end{pmatrix} = W^{(1)} \underline{x} \in \mathbb{R}^M$$

Hidden units z_m

$$z_m = h(a_m)$$

Output units y_k

$$\underline{y} = \begin{pmatrix} y_1 \\ \vdots \\ y_K \end{pmatrix} = W^{(2)} \underline{z} \in \mathbb{R}^K$$

Activation functions

$$h^{(1)}, h^{(2)}$$

Activation functions

We need activation functions!
Otherwise the NN is just
a linear model...




Figure: Popular activation functions.

Green: Logistic sigmoid

$$h(a) = \sigma(a) = \frac{1}{1 + e^{-a}}$$

Red: Hyperbolic tan

$$h(a) = \tanh(a) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

Blue:

$$h(a) = \text{ReLU}(a) = \max(0, a)$$

(Rectified Linear Unit)

Linear regression & classification as NN

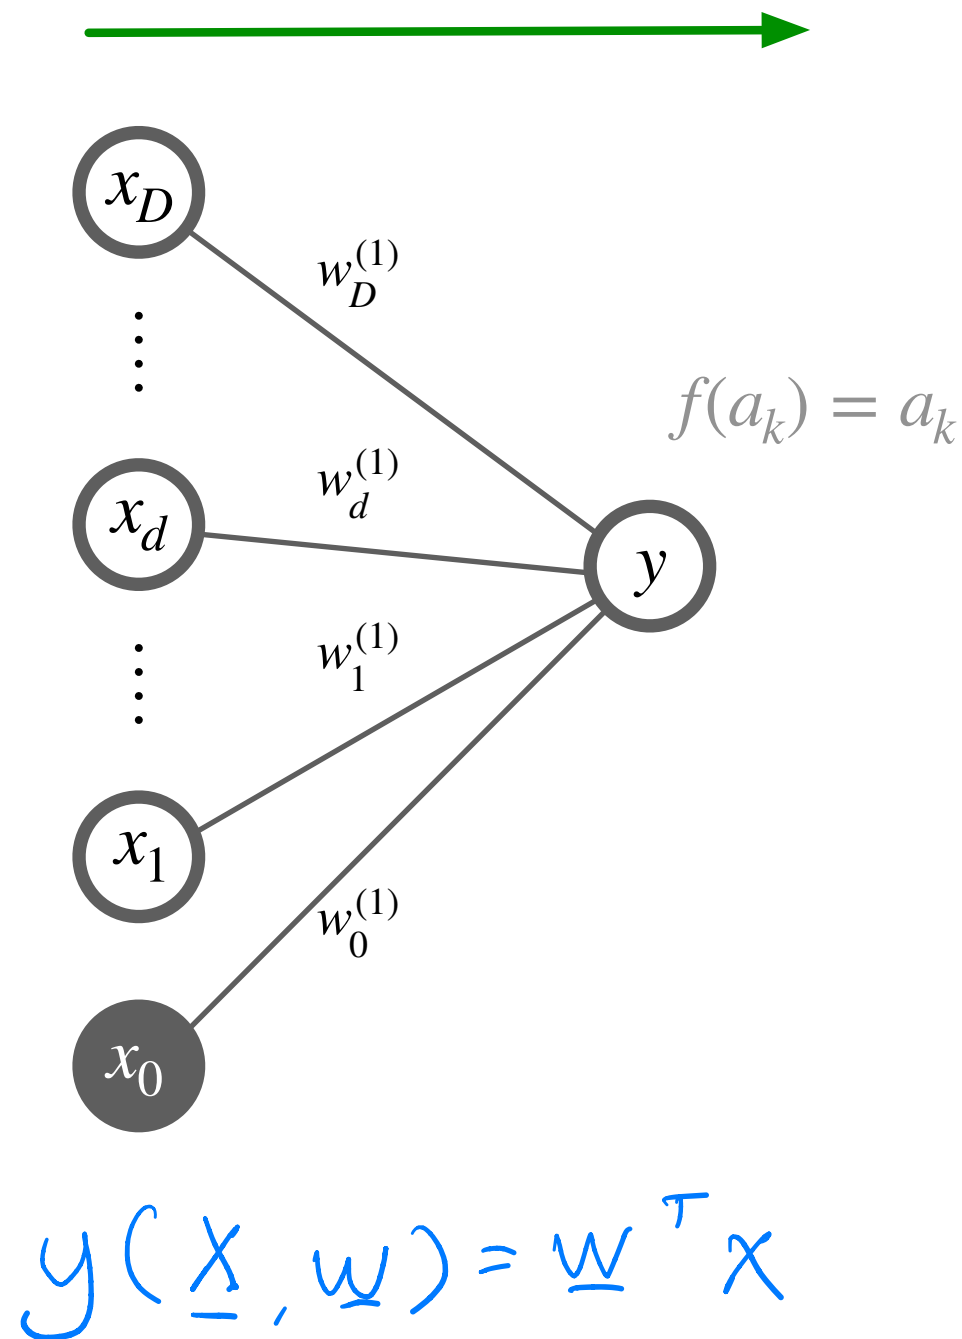


Figure: Linear regression as 1-layer NN

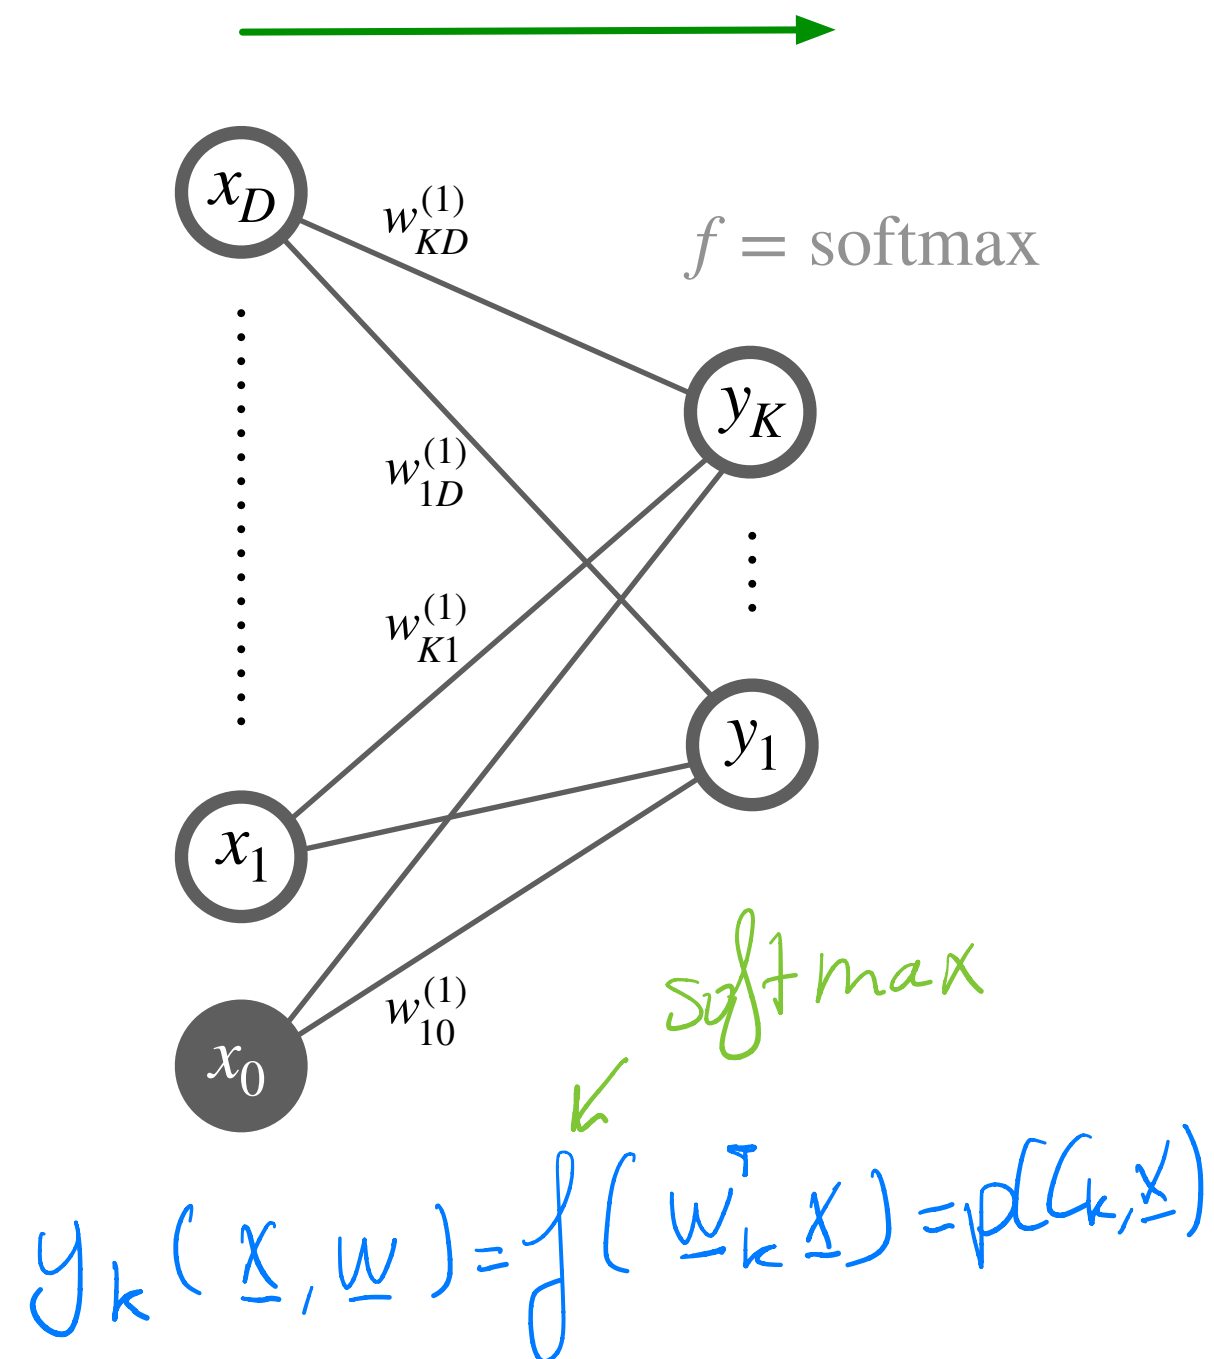
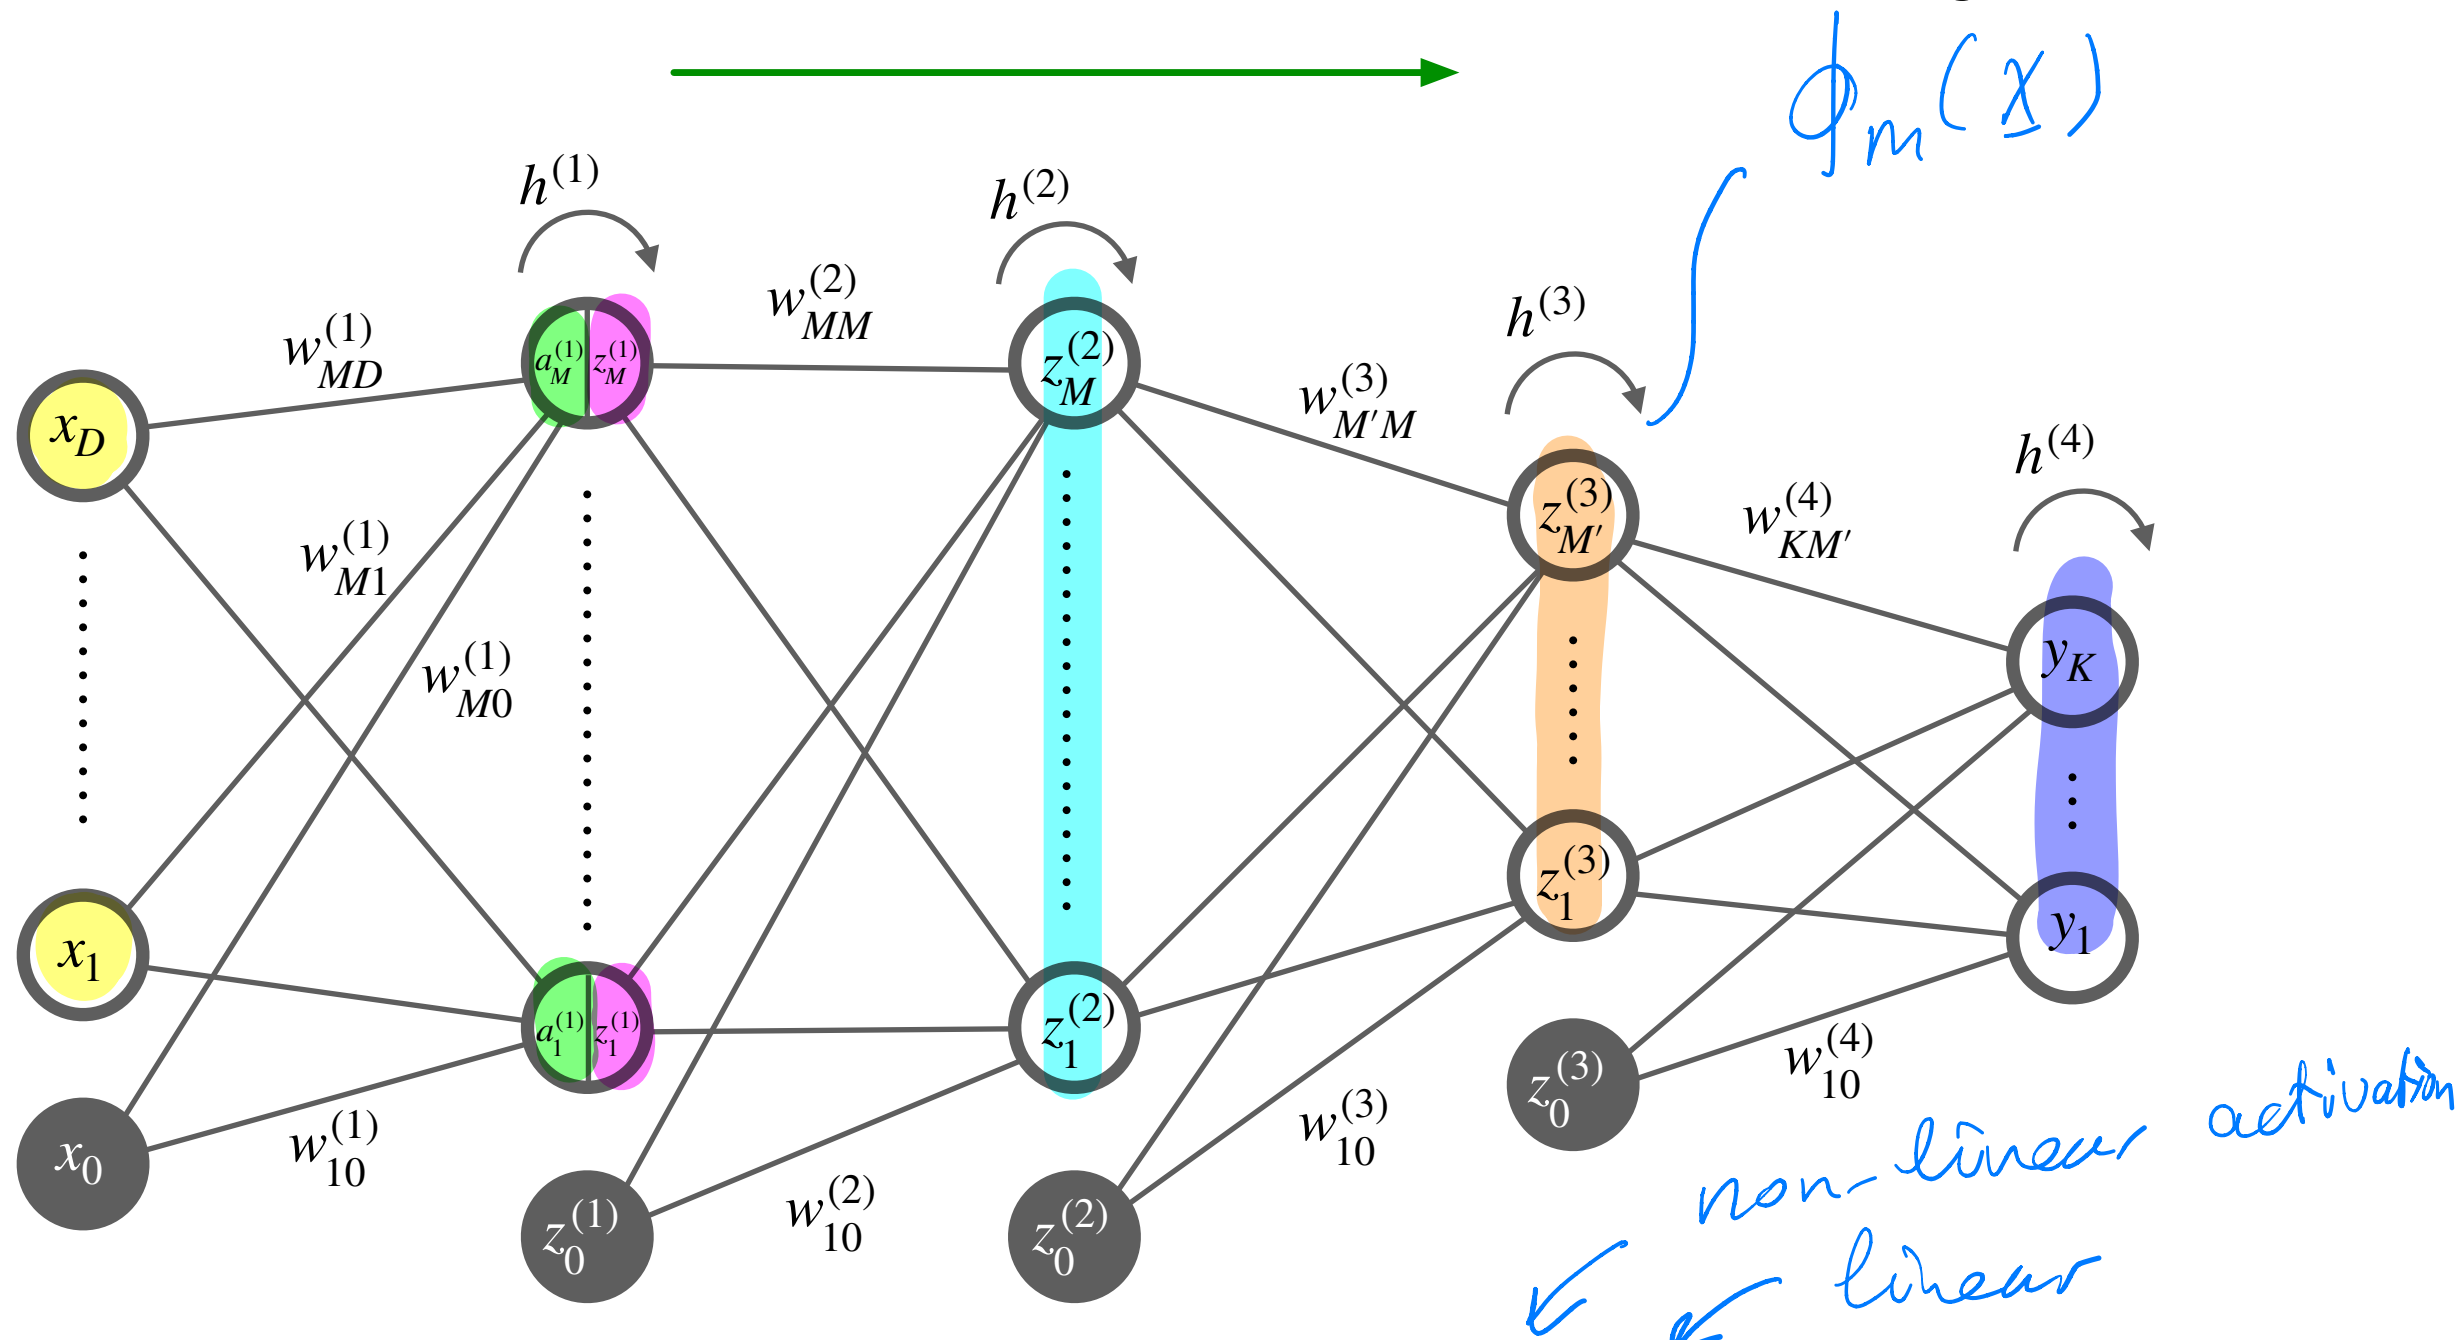


Figure: Linear Classification with K classes as 1-layer NN

Feed-Forward Networks: Multiple layers



$$y_k(\mathbf{x}, \mathbf{w}) = h^{(4)}(a^{(4)}(h^{(3)}(a^{(3)}(h^{(2)}(a^{(2)}(h^{(1)}(a^{(1)}(\mathbf{x}))))))))$$

$$= h^{(4)} \circ a^{(4)} \circ h^{(3)} \circ a^{(3)} \circ h^{(2)} \circ a^{(2)} \circ h^{(1)} \circ a^{(1)}(\mathbf{x})$$

Figure: 4 layer network. Number of layers = number of layers of adaptive weights.

Feed-Forward Networks: Skip Connections




Figure: 3 layer feed-forward net with skip connections

Feed-Forward Networks: Sparse Connections


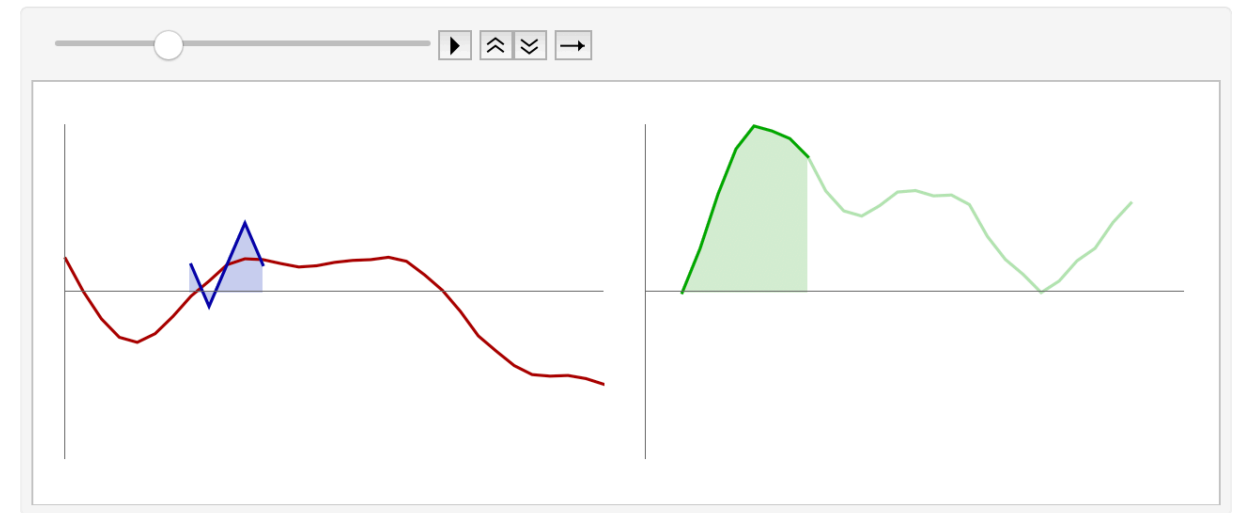



Figure: Feed-forward architecture with sparse connections. With special weight sharing --> Convolutional Neural Nets (Le Cun et al 1989)

Feed-Forward Networks: Sparse Connections

$$\underline{a} = \underline{x} \star \underline{w}$$

$$a_j = \sum_{|i-j| \leq k} x_i w_{i-j}$$



$$\underline{x} \star \underline{w} = \underline{a}$$

Example: 1D convolution = sparse + weightsharing:

$$\begin{pmatrix} a_1 \\ a_2 \\ a_2 \\ \vdots \\ a_M \end{pmatrix} = \begin{pmatrix} w_{11} & w_{12} & w_{13} & w_{14} & w_{15} & \dots \\ w_{21} & w_{22} & w_{23} & w_{24} & w_{25} & \dots \\ w_{31} & w_{32} & w_{33} & w_{34} & w_{35} & \dots \\ w_{41} & w_{42} & w_{43} & w_{44} & w_{45} & \dots \\ \vdots & \vdots & \vdots & \ddots & \ddots & \ddots \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_D \end{pmatrix}$$

sparse + weight sharing

$$= \begin{pmatrix} w_1 & w_2 & w_3 & 0 & 0 & \dots \\ 0 & w_1 & w_2 & w_3 & 0 & \dots \\ 0 & 0 & w_1 & w_2 & w_3 & \dots \\ 0 & 0 & 0 & w_1 & w_2 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_D \end{pmatrix}$$

Figure: Feed-forward architecture with sparse connections. With special weight sharing --> Convolutional Neural Nets (Le Cun et al 1989)

General Feed-Forward Architectures

- Each unit (hidden & output) in feed-forward architectures computes a function of the form

$$z_m = h \left(\sum_j w_{mj} z_j \right)$$

- No closed directed cycles!

any hidden unit
from lower layer




Figure: example of general feed-forward architecture