


Machine Learning 1

Lecture 12.2 - Kernel Methods
Gaussian Processes - Kernelizing Bayesian
Linear Regression

Erik Bekkers

(Bishop 6.4.0, 6.4.1)



Revisit Bayesian linear regression

- Given input \mathbf{x} and target t , we assumed

$$t = \boldsymbol{\phi}(\mathbf{x})^T \mathbf{w} + \varepsilon \quad \text{with} \quad \varepsilon \sim \mathcal{N}(0, \beta^{-1})$$

- So for observations $\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_N^T \end{bmatrix}$ and $\mathbf{t} = \begin{bmatrix} t_1 \\ \vdots \\ t_N \end{bmatrix}$ we have *Likelihood!*

$$\mathbf{t} = \boldsymbol{\Phi} \mathbf{w} + \boldsymbol{\varepsilon}, \quad \boldsymbol{\varepsilon} \sim \mathcal{N}(\mathbf{0}, \beta^{-1} \mathbf{I}), \quad p(\mathbf{t} | \mathbf{X}, \mathbf{w}) = \mathcal{N}(\mathbf{t} | \boldsymbol{\Phi} \mathbf{w}, \beta^{-1} \mathbf{I})$$

- We imposed a prior on \mathbf{w} :

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w} | \mathbf{0}, \boldsymbol{\Sigma}_p)$$

- After observing \mathbf{X} and \mathbf{t} , we obtained posterior

$$p(\mathbf{w} | \mathbf{X}, \mathbf{t}) = \frac{p(\mathbf{t} | \mathbf{X}, \mathbf{w}) p(\mathbf{w})}{p(\mathbf{t} | \mathbf{X})} = \mathcal{N}(\mathbf{w} | \mathbf{m}_N, \mathbf{S}_N)$$

$$\text{with} \quad \mathbf{m}_N = \beta \mathbf{S}_N^{-1} \boldsymbol{\Phi}^T \mathbf{t}, \quad \mathbf{S}_N^{-1} = \beta \boldsymbol{\Phi}^T \boldsymbol{\Phi} + \boldsymbol{\Sigma}_p^{-1}$$

Revisit Bayesian linear regression

- Predictions for new point \mathbf{x}^* , after observing N data points

$$p(t^* | \mathbf{x}^*, \mathbf{X}, \mathbf{t}) = \int p(t^* | \mathbf{x}^*, \mathbf{w}) p(\mathbf{w} | \mathbf{X}, \mathbf{t}) d\mathbf{w}$$

$$= \mathcal{N}(t^* | \mu_N(\mathbf{x}^*), \sigma_N^2(\mathbf{x}^*))$$





with

$$\mu_N(\mathbf{x}^*) = \boldsymbol{\phi}(\mathbf{x}^*)^T \mathbf{m}_N = \beta \boldsymbol{\phi}(\mathbf{x}^*)^T \mathbf{S}_N^{-1} \boldsymbol{\Phi}^T \mathbf{t} = \sum_{n=1}^N \underbrace{\beta \boldsymbol{\phi}(\mathbf{x}^*)^T \mathbf{S}_N^{-1} \boldsymbol{\phi}(\mathbf{x}_n)}_{k(\mathbf{x}^*, \mathbf{x}_n)} t_n$$

$$\sigma_N^2(\mathbf{x}^*) = \beta^{-1} + \boldsymbol{\phi}(\mathbf{x}^*)^T \mathbf{S}_N \boldsymbol{\phi}(\mathbf{x}^*)$$

- Mean predictions are linear combinations of kernels!

Revisit Bayesian linear regression



Revisit Bayesian linear regression

- ▶ Linear Bayesian regression suffers from limited expressiveness
(linear in w)
- ▶ Can become more expressive for larger M , where $\phi(\mathbf{x}) \in \mathbb{R}^M$
- ▶ However, this also requires more parameters! ($\mathbf{w} \in \mathbb{R}^M$)
- ▶ It's expensive to train more expressive models! (requires inverting \mathbf{S}_N which is of size $M \times M$)
- ▶ We also need to choose our basis functions at 'good' locations
- ▶ But we saw that the final mean prediction can be written in equivalent kernel form...

